

```

from google.colab import drive
import numpy as np
import cv2, os, torch
import torch.nn as nn
from torch.utils.data import DataLoader, Dataset
from sklearn.metrics import accuracy_score

```

## ▼ Read the images (function from the first seminar)

```
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount(force=True)



```

def read_files(X, Y, path, ans):
    files = os.listdir(path)
    for name in files:
        img = cv2.imread(path + '/' + name, 0)
        if img.shape != 0:
            img = cv2.resize(img, (256, 256))
            vect = img.reshape(1, 256 * 2)
            vect = vect / 255.
            X = vect if (X is None) else np.vstack((X, vect))
            Y = np.append(Y, ans)
    return X, Y

```

```

path = "/content/drive/MyDrive/lesson1_dataset"
X = None
Y = np.array([])
X, Y = read_files(X, Y, path + "/logloss_0", 0)
X, Y = read_files(X, Y, path + "/logloss_1", 1)

```

## ▼ Create the dataset

```

class CreateDataset(Dataset):
    def __init__(self, X, Y):
        self.x = torch.from_numpy(X)
        self.y = torch.from_numpy(Y)

    def __getitem__(self, idx):
        return self.x[idx], self.y[idx]

    def __len__(self):
        return self.x.shape[0]

```

```
np.random.seed(1)
```

```

np.random.seed(1)
dataset = CreateDataset(X, Y)
train, test = torch.utils.data.random_split(dataset, [round(0.8*len(dataset)), round(0.2*len(dataset))])

train_loader = DataLoader(dataset=train, batch_size=16, shuffle=True)
test_loader = DataLoader(dataset=test, shuffle=True)

for x,y in train_loader:
    print(x.view(x.shape[0], -1).shape, y.shape)
    break

    torch.Size([16, 65536]) torch.Size([16])

```

## ▼ Configure the model

```

model = nn.Sequential(
    nn.Linear(65536, 2048),
    nn.BatchNorm1d(2048),
    nn.Dropout(0.5),
    nn.ReLU(),
    nn.Linear(2048, 2048),
    nn.BatchNorm1d(2048),
    nn.ReLU(),
    nn.Linear(2048, 2048),
    nn.BatchNorm1d(2048),
    nn.ReLU(),
    nn.Linear(2048, 2048),
    nn.BatchNorm1d(2048),
    nn.Dropout(0.5),
    nn.ReLU(),
    nn.Linear(2048, 2),
    nn.Softmax()
)

criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters())

model.cuda()

Sequential(
  (0): Linear(in_features=65536, out_features=2048, bias=True)
  (1): BatchNorm1d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (2): Dropout(p=0.5, inplace=False)
  (3): ReLU()
  (4): Linear(in_features=2048, out_features=2048, bias=True)
  (5): BatchNorm1d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (6): ReLU()
  (7): Linear(in_features=2048, out_features=2048, bias=True)
  (8): BatchNorm1d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (9): ReLU()
  (10): Linear(in_features=2048, out_features=2048, bias=True)

```

```

(11): BatchNorm1d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(12): Dropout(p=0.5, inplace=False)
(13): ReLU()
(14): Linear(in_features=2048, out_features=2, bias=True)
(15): Softmax(dim=None)
)

```

## ▼ Training

```

epochs = 20
model.train()
for i in range(epochs):
    for j, (x, y) in enumerate(train_loader):
        optimizer.zero_grad()
        x = x.view(x.shape[0], -1)
        x = x.cuda()
        y = y.cuda()
        y_pred = model(x.float())
        loss = criterion(y_pred, y.long())
        print(f"Epoch {i}\t iter {j}\t loss {loss}")
        loss.backward()
        optimizer.step()

```

```

Epoch 1  iter 0  loss 0.5469618439674377
Epoch 1  iter 1  loss 0.3879917562007904
Epoch 1  iter 2  loss 0.4894407391548157
Epoch 1  iter 3  loss 0.31327298283576965
Epoch 2  iter 0  loss 0.49811315536499023
Epoch 2  iter 1  loss 0.3730299472808838
Epoch 2  iter 2  loss 0.37833499908447266
Epoch 2  iter 3  loss 0.5633803606033325
Epoch 3  iter 0  loss 0.4455563426017761
Epoch 3  iter 1  loss 0.47449222207069397
Epoch 3  iter 2  loss 0.4343473017215729
Epoch 3  iter 3  loss 0.5711500644683838
Epoch 4  iter 0  loss 0.3688523471355438
Epoch 4  iter 1  loss 0.37728050351142883
Epoch 4  iter 2  loss 0.46270787715911865
Epoch 4  iter 3  loss 0.5044106245040894
Epoch 5  iter 0  loss 0.4518081545829773
Epoch 5  iter 1  loss 0.3159264922142029
Epoch 5  iter 2  loss 0.3764341175556183
Epoch 5  iter 3  loss 0.7146843671798706
Epoch 6  iter 0  loss 0.39910557866096497
Epoch 6  iter 1  loss 0.3561722934246063
Epoch 6  iter 2  loss 0.3683589994907379
Epoch 6  iter 3  loss 0.5638160705566406
Epoch 7  iter 0  loss 0.31502851843833923

Epoch 7  iter 1  loss 0.4620009660720825
Epoch 7  iter 2  loss 0.33126240968704224
Epoch 7  iter 3  loss 0.3868955969810486
Epoch 8  iter 0  loss 0.4795559346675873
Epoch 8  iter 1  loss 0.34585002064704895
Epoch 8  iter 2  loss 0.36920568346977234
Epoch 8  iter 3  loss 0.5638471374511710

```

```

Epoch 8   iter 3   loss 0.5628471374511719
Epoch 9   iter 0   loss 0.3749336898326874
Epoch 9   iter 1   loss 0.3132951557636261
Epoch 9   iter 2   loss 0.4352795481681824
Epoch 9   iter 3   loss 0.3132992386817932
Epoch 10   iter 0   loss 0.31454357504844666
Epoch 10   iter 1   loss 0.4150834381580353
Epoch 10   iter 2   loss 0.4095011353492737
Epoch 10   iter 3   loss 0.5631428360939026
Epoch 11   iter 0   loss 0.37777575850486755
Epoch 11   iter 1   loss 0.5640539526939392
Epoch 11   iter 2   loss 0.3147273063659668
Epoch 11   iter 3   loss 0.5641259551048279
Epoch 12   iter 0   loss 0.41222891211509705
Epoch 12   iter 1   loss 0.5009142756462097
Epoch 12   iter 2   loss 0.3133174479007721
Epoch 12   iter 3   loss 0.31327205896377563
Epoch 13   iter 0   loss 0.31412893533706665
Epoch 13   iter 1   loss 0.31390100717544556
Epoch 13   iter 2   loss 0.4682944715023041
Epoch 13   iter 3   loss 0.8066140413284302
Epoch 14   iter 0   loss 0.31338104605674744
Epoch 14   iter 1   loss 0.4227094054222107
Epoch 14   iter 2   loss 0.3791202902793884
Epoch 14   iter 3   loss 0.3463093340396881
Epoch 15   iter 0   loss 0.376575767993927
Epoch 15   iter 1   loss 0.313591867685318

```

## ▼ Evaluation of the model

```

model.eval()
y_true = []; y_pred = []
for x, y in test_loader:
    x = x.cuda()
    y = y.cuda()
    y_pred.append(round(model(x.float()).data[0][1].item()))
    y_true.append(int(y))

```

```
print('The accuracy of the model is:', accuracy_score(y_true, y_pred))
```

```

The accuracy of the model is: 0.9230769230769231
/usr/local/lib/python3.6/dist-packages/torch/nn/modules/container.py:117: UserWarning
    input = module(input)

```